

## Known Issues :

First and foremost, I want to point out some strange unexpected behaviour that I ran into from the program. There is a function called `selectHyperParams` which depending on the function name that we pass to it, does cross validation over 5 passes, and one epoch, with that function.

After choosing the parameter, I run the training experiment with that parameter as the passed hyper parameter.

For some reason, **simply calling the `selectHyperParams` function changes the result shown by the second call to the training perceptron.** Even when I'm not passing the hyper parameter selected by the previous function. So while **without running the function, it shows number of mistakes to be around 60, after calling the function it shows number of mistakes as about 260.**

ALL my variables other than `numRuns` have been defined inside the context of a function.

As a work around, I have commented the `selectHyperParams` function in the second part of the experiment (Although, thankfully as per the language, it felt more like we could try out a few values and plug in what we thought best)

I also did the experiment while commenting out this particular function in the rest of the code with hard coded hyper parameters, and I think I got better results. If you have any insight, do let me know.

## Experiment 1: Sanity Check

The sanity check works out. The weight vector shows the most significant attribute as **X4 with a weight of more than 6** for a rate of 3. Far more than any of the other vectors.

## Experiment 2: One Pass for Normal & Margin Perceptron

Three functions are used here.

1. `trainPerceptron()` : Accuracy on **test** set : **1** **training** set : **0.969**
2. `trainMarginPerceptron()`: Accuracy on **test** set : **0.999** **training** set : **0.912**
3. `testAnyPerceptron()` : Tests the perceptron. Takes weight and records as params

For both, by default, if no weight vector is passed, the function generates it randomly. However a particular initial weight vector can also be passed to the function. I tried a number of values for the hyper parameters, and left the ones that worked best. An issue I came across has been mentioned above.

As per the algorithm of the perceptron, both the functions call 2 methods the `dotProd` which returns the dot product of the 2 passed vectors (`Xvector` and `Weight vector`), as well as the weight vector. The weight vector can be updated using the `updateWvec` function which multiplies the `Xvector` with the label, and rate value and adds it to the old weight vector value.

While for the perceptron, this change is made when  $y \cdot \text{dotProduct} < 0$ , for margin it is for  $< \mu$  where  $\mu$  is a margin value.

## Experiment 3: 10 Epochs for all data sets

Five functions are used here.

1. `trainPerceptron()`
2. `trainMarginPerceptron()`
3. `trainAggressiveMarginPerceptron()`
4. `testAnyPerceptron`
5. `selectHyperParams()`

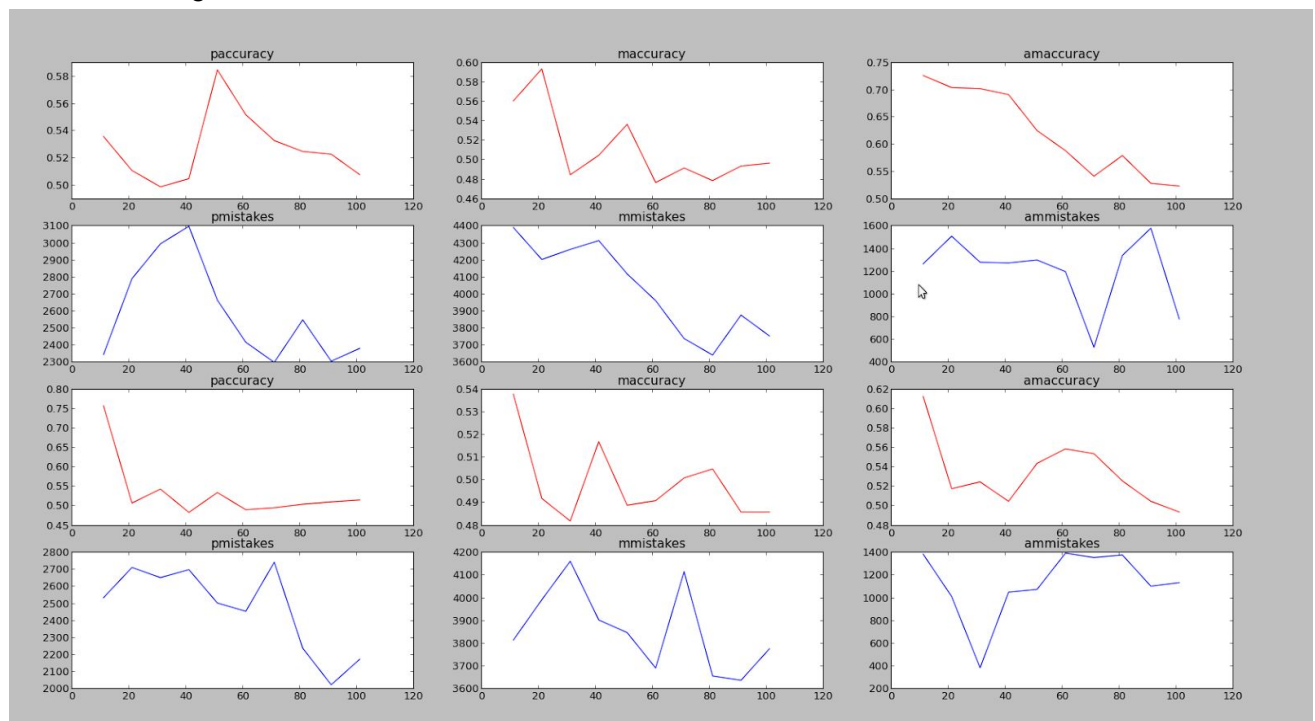
The functioning of the first 2 perceptrons is already described above. For the aggressive perceptron, rate is not a passed parameter but is calculated based on the formula in the assignment. The behaviour of functions 4 and 5 is described above.

This experiment is carried out over 10 epochs from each function, coupled with a 5 fold validation while selecting the hyper parameters.

I have also tried to make the listing out of results a readable process, and the program produces graphs at the end of the computation with any values you wish to tweak.

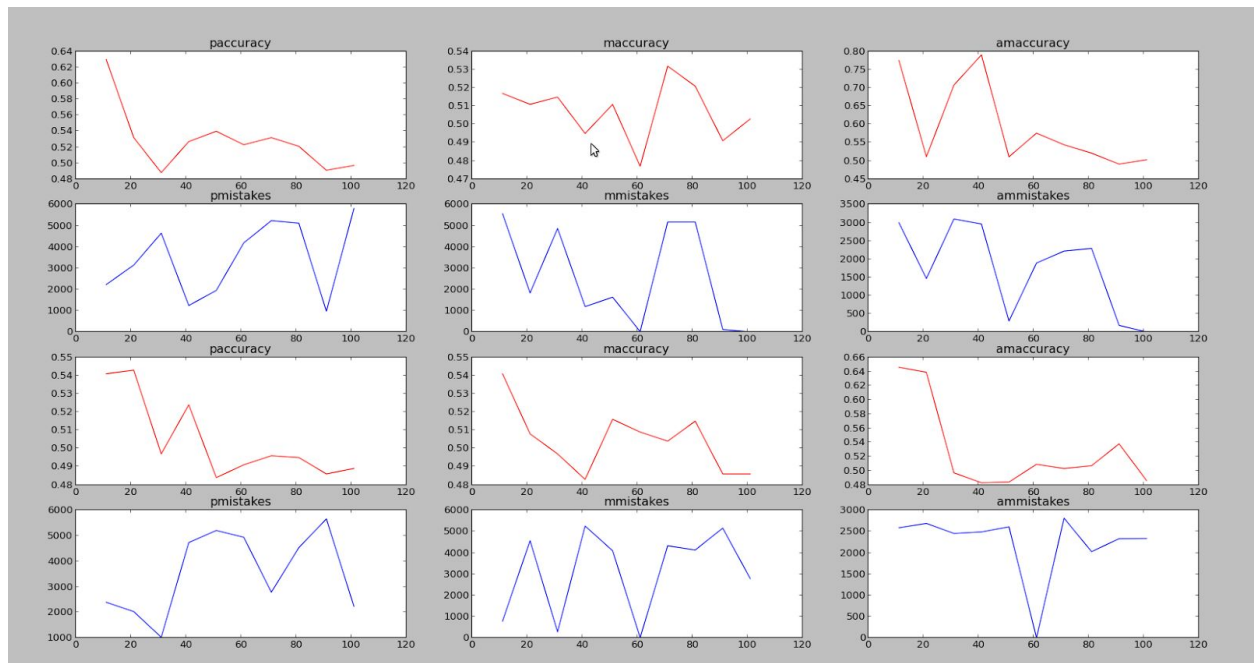
## Result

It is quite clear that while the accuracy of the regular perceptron decreases with increasing dimensionality, the margin and aggressive margin perceptron perform better in comparison. The first 6 images are from data0, and next 6 from data1



The below was the recorded performance from the same experiment, with a few tweaks.

One of which was switching on the selectHyperParams function. The data seems consistent save a few outliers with low accuracy.



The reason for stating those issues was more out of curiosity if you have any ideas about what could be causing that. I suppose in retrospect they don't affect the program that badly since the later data seems to be fairly consistent with expected results save a few outliers.